

The Analysis of Fuel Prices in the city of Ahmedabad

Name: Juilee Sameer Thakur

Date: 2023-05-07

Name of the course: MA 641 - Time Series Analysis

Introduction

Increasing fuel prices have been a constant concern in the city of Ahmedabad, India. In this project, I analysed the daily fuel data using time series analysis methods to better understand the data and its aspects and thus help predict the future prices of fuel.

Importing the dataset

```
price_data <- read.csv("~/Stevens Institute of Technology/Spring 2023/MA 641/Project/Non-seasonal/Ahmedabad.csv")
head(price_data)
```

```
##      city      date  rate
## 1 Ahmedabad 2019-04-22 70.32
## 2 Ahmedabad 2019-04-22 70.32
## 3 Ahmedabad 2019-04-22 70.32
## 4 Ahmedabad 2019-04-22 70.32
## 5 Ahmedabad 2019-04-22 70.32
## 6 Ahmedabad 2019-04-22 70.32
```

The data-set contains information of fuel rate of last 4 years in Ahmedabad, India. That data set contains 1024 observations and has 2 major attributes.

- date - This column contains the date at which the rate was recorded.
- rate - This column contains the rate of the fuel on the corresponding date.

Summarizing the data-set

Below, I summarized the data to see the values of each attribute of the data-set and check for any missing values in the data-set.

```
summary(price_data)
```

```
##      city      date      rate
## Length:1024      Length:1024      Min.   : 67.15
## Class :character      Class :character      1st Qu.: 70.39
## Mode  :character      Mode  :character      Median : 70.81
##                                     Mean    : 79.45
##                                     3rd Qu.: 89.00
##                                     Max.    :106.63
```

The average fuel rate is Rs. 79.45 with maximum rate being Rs 106.63. I found out that the data-set has no missing values. Hence, we can now proceed with our analysis. I am keen on knowing the future prices of fuel hence, the desired attribute of study is 'rate'.

Checking for seasonality

To check if the pricedata is seasonal or not, I run the combined test from the "seastests" library. The combined test can be used individually or using the *isSeasonal()* function. So, I ran both tests.

```
library(seastests)
```

```
## Warning: package 'seastests' was built under R version 4.2.3
```

```
# Combined test
combined_test(price_data$rate, freq = 12)
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
## as.zoo.data.frame zoo
```

```
## Test used:  W0
##
## Test statistic:  0
## P-value:  1 1 0.9374157
```

```
# Seasonal or not Test
isSeasonal(price_data$rate, test = "combined", freq = 12)
```

```
## [1] FALSE
```

Since overall p-value of combined test is 0.937 which is considerably close to 1 than 0, we say the data is non-seasonal. Just to get a confirmation, I also applied the isSeasonal test which returned false. Thus, concluded the data-set is Non-seasonal.

Exploratory Data Analysis

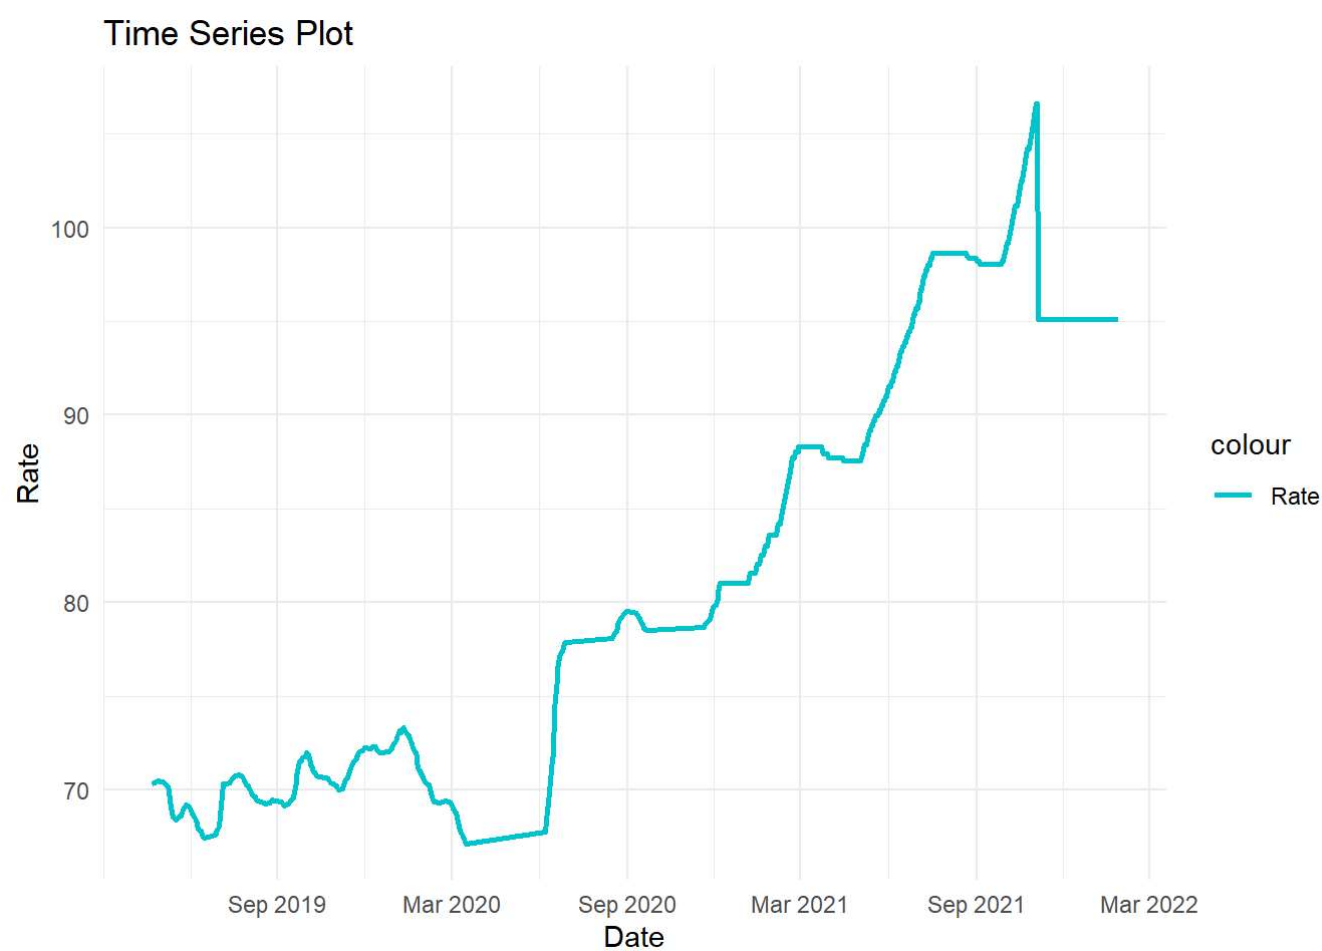
I performed some exploratory data analysis to check out the time series plot and see the trend in data over time.

```
options(warn = -1)

library(ggplot2)

price_data$date <- as.Date(price_data$date)

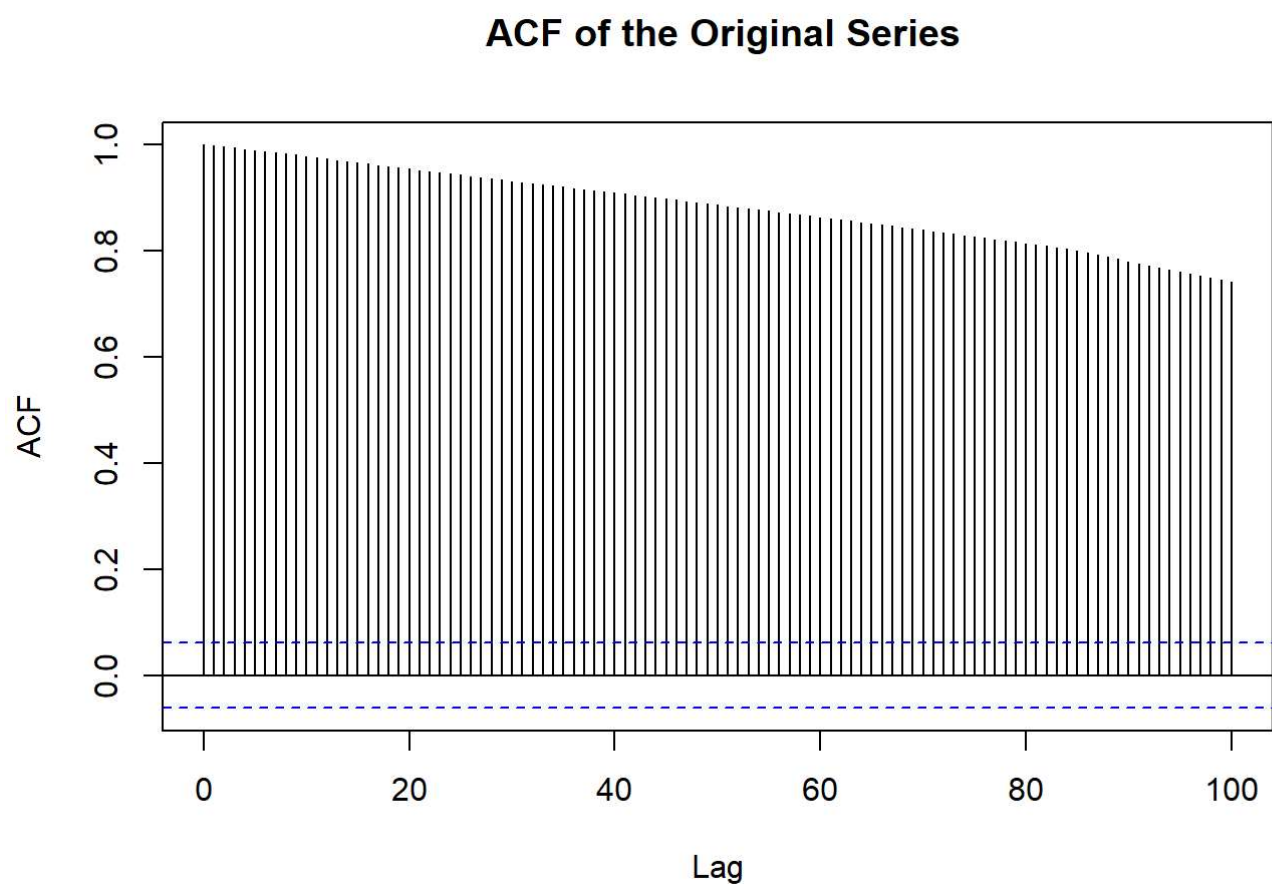
ggplot(price_data, aes(x = date, group = 1)) +
  geom_line(aes(y = rate, color = "Rate"), linewidth = 1) +
  labs(x = "Date", y = "Rate", title = "Time Series Plot") +
  scale_x_date(date_breaks = "6 month", date_labels = "%b %Y") +
  scale_color_manual(values = c("Rate" = "turquoise3")) +
  theme_minimal()
```



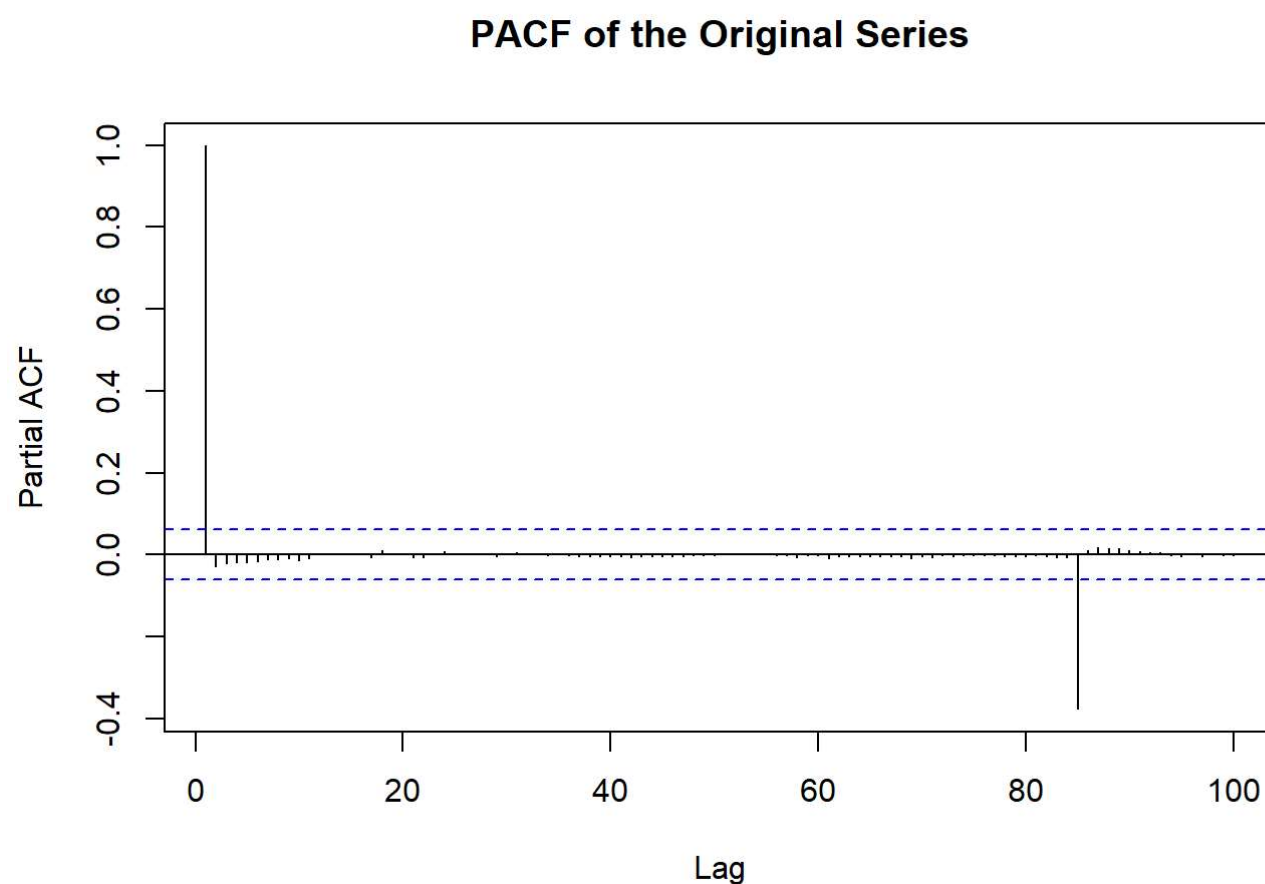
The data seems to follow an upward trend which makes sense as it is the price data for fuels. We now check for the ACF and PACF of the series.

ACF and PACF of the series

```
acf(price_data$rate, lag.max = 100, main = " ACF of the Original Series")
```



```
pacf(price_data$rate, lag.max = 100, main = "PACF of the Original Series")
```



The data shows non-seasonal characteristics but in order to perform further analysis we check for the stationarity of the series.

Checking the stationarity of the series

In order to check the stationarity, I performed the Augmented Dicky-Fuller test from the 'tseries' library.

Following the hypothesis for ADF test:

H_0 : The series is non-stationary

vs

H_1 : The series is stationary

```
options(warn = -1)

library(tseries)
adf.test(price_data$rate)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: price_data$rate
## Dickey-Fuller = -2.011, Lag order = 10, p-value = 0.5737
## alternative hypothesis: stationary
```

Since p-value is greater than 0.05, we accept H_0 and say that the data is not stationary. In order to make the data stationary, I initially took the difference of the data but to make analysis better I took the difference of the log of the series and made it stationary.

```
library(tseries)
diff_price = diff(log(price_data$rate))
adf.test(diff_price)
```

```
##
## Augmented Dickey-Fuller Test
##
## data: diff_price
## Dickey-Fuller = -8.1548, Lag order = 10, p-value = 0.01
## alternative hypothesis: stationary
```


After transforming the data, I got the Augmented Dicky-Fuller test results which gave the p-value as 0.01 which in reality was much smaller. Hence, we can say that taking the difference of the logarithm of the data made it stationary. We shall now use this transformed series to perform further analysis.

Finding the model

Plotting ACF and PACF of the transformed series.

```
acf(diff_price, lag.max = 50, main = "ACF of the Transformed Series")
```


ACF of the Transformed Series



We can observe 5 significant lags here.

```
pacf(diff_price, lag.max = 50, main = "PACF of the Transformed Series")
```

PACF of the Transformed Series



We can observe 4 significant lags.

Model can be suggested as ARIMA(5,1,4) as we did one differencing so $d = 1$.

```
library(TSA)
```

```
## Registered S3 methods overwritten by 'TSA':  
##   method      from  
## fitted.Arima forecast  
## plot.Arima   forecast
```

```
##  
## Attaching package: 'TSA'
```

```
## The following objects are masked from 'package:stats':  
##  
##   acf, arima
```

```
## The following object is masked from 'package:utils':  
##  
##   tar
```

```
eacf(diff_price)
```

```
## AR/MA
##   0 1 2 3 4 5 6 7 8 9 10 11 12 13
## 0 x x x x x o o o o o o o o
## 1 x o o o o o o o o o o o o
## 2 x x o o o o o o o o o o o
## 3 x x x o o o o o o o o o
## 4 x x x x o o o o o o o o
## 5 x x o o x o o o o o o o
## 6 x x x o x x o o o o o o
## 7 x x x x o x x o o o o o
```

We can deduce from the EACF that the model can be ARMA(1,1,1), ARIMA(1,1,2) or ARIMA(2,1,2).

Thus models suggested are:

- Model 1: ARIMA (5,1,4)
- Model 2: ARIMA (1,1,1)
- Model 3: ARIMA(1,1,2)
- Model 4: ARIMA (2,1,2)

Parameter Estimation

Now, for each of the above models, we check the AIC and BIC values.

```
library(forecast)
model_1 <- Arima(diff_price, order = c(5,1,4))
model_2 <- Arima(diff_price, order = c(1,1,1))
model_3 <- Arima(diff_price, order = c(1,1,2))
model_4 <- Arima(diff_price, order = c(2,1,2))
#summary(model_1)
```

```
cat("For Model 1:\n")
```

```
## For Model 1:
```

```
print(paste("AIC is", model_1$bic))
```

```
## [1] "AIC is -8250.82488710856"
```

```
print(paste("BIC is", model_1$bic))
```

```
## [1] "BIC is -8250.82488710856"
```

```
cat("\nFor Model 2:\n")
```

```
##
## For Model 2:
```

```
print(paste("AIC is", model_2$aic))
```

```
## [1] "AIC is -8292.82927131186"
```

```
print(paste("BIC is", model_2$bic))
```

```
## [1] "BIC is -8278.04072099957"
```

```
cat("\nFor Model 3:\n")
```

```
##
## For Model 3:
```

```
print(paste("AIC is", model_3$aic))
```

```
## [1] "AIC is -8311.16130092161"
```

```
print(paste("BIC is", model_3$bic))
```

```
## [1] "BIC is -8291.44323383855"
```

```
cat("\nFor Model 4:\n")
```

```
##  
## For Model 4:
```

```
print(paste("AIC is", model_4$aic))
```

```
## [1] "AIC is -8309.35938509031"
```

```
print(paste("BIC is", model_4$bic))
```

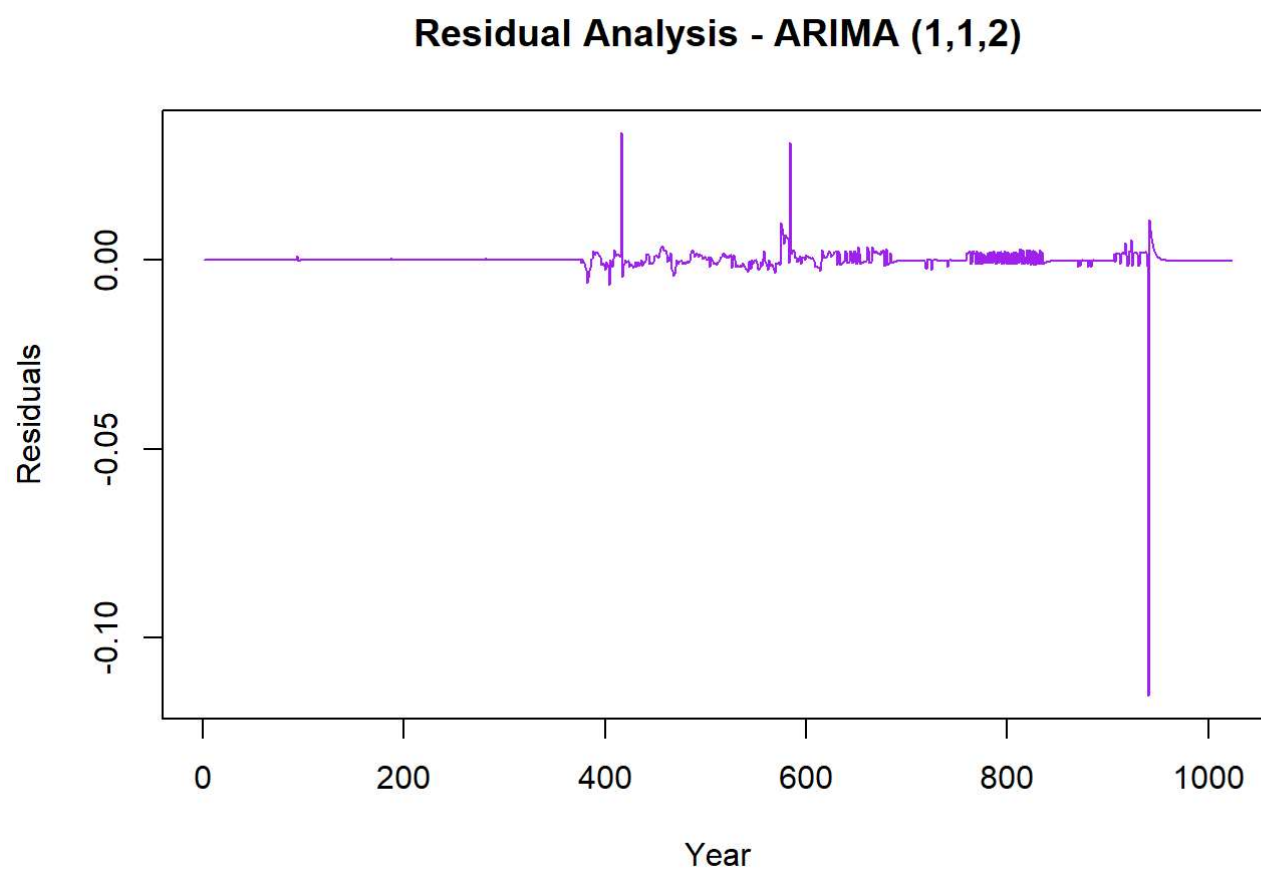
```
## [1] "BIC is -8284.71180123649"
```

The lowest AIC and BIC values are for Model 3 i.e. ARIMA (1,1,2).

Residual Analysis

Using the above obtained model, I performed the residual analysis and plotted the same along with its ACF and PACF.


```
ts.plot(model_3$residuals,lwd=1,col="purple",xlab = "Year", ylab = "Residuals",main='Residual Analysis - ARIMA (1,1,2)')
```



```
#plot(residu, xlab = "Year", ylab = "Residuals")
```


```
acf(model_3$residuals, lag.max = 100, main = "ACF of the Residuals")
```

ACF of the Residuals



```
pacf(model_3$residuals, lag.max = 100, main = "PACF of the Residuals")
```


PACF of the Residuals



The ACF and the PACF plots portray white noise hence indicating ideal for forecasting.


Checking for normality

```
qqnorm(model_3$residuals)  
qqline(model_3$residuals)
```



The q-q plot shows partial normality as towards the either ends, the plot seems to distort around the q-q line.

```
hist(model_3$residuals, main = "Histogram of Residuals")
```



```
shapiro.test(model_3$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model_3$residuals
## W = 0.18304, p-value < 2.2e-16
```

The p-value is significantly less than 0.05. Hence, the conclusion is the distribution is non-normal.

```
Box.test(model_3$residuals, type = "Ljung-Box")
```

```
##
##  Box-Ljung test
##
## data:  model_3$residuals
## X-squared = 0.066823, df = 1, p-value = 0.796
```

Thus, we fail to reject the null hypothesis of the test and conclude that the data values are independent.

Forecasting

Using the above fitted model ARIMA(1,1,2), we plot the forecasted values on original plot.


```
library(forecast)
d <- Arima(price_data$rate, order = c(1,1,2))
forecasted_series <- forecast(d, h = 90)

cat("The confidence interval is", forecasted_series$level)
```

```
## The confidence interval is 80 95
```


```
plot(forecasted_series)
```

Forecasts from ARIMA(1,1,2)



```
plot(forecasted_series, PI = TRUE, include = 80)
```

Forecasts from ARIMA(1,1,2)



```
cat("\nFollowing is the summary of the forecasted series:")
```

```
##
## Following is the summary of the forecasted series:
```

```
summary(forecasted_series)
```

```
##
## Forecast method: ARIMA(1,1,2)
##
## Model Information:
## Series: price_data$rate
## ARIMA(1,1,2)
##
## Coefficients:
##      ar1      ma1      ma2
##      0.8366 -0.765 -0.0006
## s.e.  0.0824  0.088  0.0348
##
## sigma^2 = 0.1559: log likelihood = -499.43
## AIC=1006.87  AICc=1006.91  BIC=1026.59
##
## Error measures:
##      ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 0.01694996 0.3940701 0.07540321 0.02022817 0.09151361 0.9435778
##      ACF1
## Training set -0.00176122
##
## Forecasts:
##      Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## 1025      95.11 94.60399 95.61601 94.33612 95.88388
## 1026      95.11 94.36835 95.85165 93.97574 96.24426
## 1027      95.11 94.17327 96.04673 93.67740 96.54260
## 1028      95.11 93.99906 96.22094 93.41096 96.80904
## 1029      95.11 93.83861 96.38139 93.16558 97.05442
## 1030      95.11 93.68850 96.53150 92.93600 97.28400
## 1031      95.11 93.54673 96.67327 92.71919 97.50081
## 1032      95.11 93.41204 96.80796 92.51320 97.70680
## 1033      95.11 93.28351 96.93649 92.31663 97.90337
## 1034      95.11 93.16047 97.05953 92.12844 98.09156
## 1035      95.11 93.04235 97.17765 91.94780 98.27220
## 1036      95.11 92.92872 97.29128 91.77402 98.44598
## 1037      95.11 92.81918 97.40082 91.60650 98.61350
## 1038      95.11 92.71342 97.50658 91.44475 98.77525
## 1039      95.11 92.61114 97.60886 91.28832 98.93168
## 1040      95.11 92.51208 97.70792 91.13682 99.08318
## 1041      95.11 92.41600 97.80400 90.98989 99.23011
## 1042      95.11 92.32271 97.89729 90.84721 99.37279
## 1043      95.11 92.23201 97.98799 90.70850 99.51150
## 1044      95.11 92.14374 98.07626 90.57349 99.64651
## 1045      95.11 92.05773 98.16227 90.44195 99.77805
## 1046      95.11 91.97385 98.24615 90.31367 99.90633
## 1047      95.11 91.89196 98.32804 90.18844 100.03156
## 1048      95.11 91.81196 98.40804 90.06608 100.15392
## 1049      95.11 91.73372 98.48628 89.94642 100.27358
## 1050      95.11 91.65715 98.56285 89.82932 100.39068
## 1051      95.11 91.58215 98.63785 89.71463 100.50537
## 1052      95.11 91.50865 98.71135 89.60222 100.61778
## 1053      95.11 91.43657 98.78343 89.49197 100.72803
## 1054      95.11 91.36582 98.85418 89.38377 100.83623
## 1055      95.11 91.29635 98.92365 89.27752 100.94248
## 1056      95.11 91.22808 98.99192 89.17312 101.04688
## 1057      95.11 91.16098 99.05902 89.07049 101.14951
## 1058      95.11 91.09497 99.12503 88.96954 101.25046
## 1059      95.11 91.03001 99.18999 88.87019 101.34981
## 1060      95.11 90.96605 99.25395 88.77238 101.44762
## 1061      95.11 90.90306 99.31694 88.67604 101.54396
## 1062      95.11 90.84099 99.37901 88.58111 101.63889
## 1063      95.11 90.77979 99.44021 88.48752 101.73248
## 1064      95.11 90.71945 99.50055 88.39523 101.82477
## 1065      95.11 90.65991 99.56009 88.30418 101.91582
## 1066      95.11 90.60116 99.61884 88.21432 102.00568
## 1067      95.11 90.54316 99.67684 88.12562 102.09438
## 1068      95.11 90.48588 99.73412 88.03802 102.18198
## 1069      95.11 90.42930 99.79070 87.95149 102.26851
## 1070      95.11 90.37340 99.84660 87.86599 102.35401
## 1071      95.11 90.31814 99.90186 87.78149 102.43851
## 1072      95.11 90.26352 99.95648 87.69795 102.52205
## 1073      95.11 90.20950 100.01050 87.61533 102.60467
## 1074      95.11 90.15607 100.06393 87.53362 102.68638
## 1075      95.11 90.10321 100.11679 87.45277 102.76723
## 1076      95.11 90.05090 100.16910 87.37277 102.84723
## 1077      95.11 89.99913 100.22087 87.29359 102.92641
## 1078      95.11 89.94787 100.27213 87.21520 103.00480
## 1079      95.11 89.89712 100.32288 87.13758 103.08242
## 1080      95.11 89.84685 100.37315 87.06071 103.15929
## 1081      95.11 89.79707 100.42293 86.98457 103.23543
```

## 1082	95.11	89.74774	100.47226	86.90913	103.31087
## 1083	95.11	89.69886	100.52114	86.83438	103.38562
## 1084	95.11	89.65043	100.56957	86.76030	103.45970
## 1085	95.11	89.60241	100.61759	86.68687	103.53313
## 1086	95.11	89.55481	100.66519	86.61408	103.60592
## 1087	95.11	89.50762	100.71238	86.54190	103.67810
## 1088	95.11	89.46082	100.75918	86.47033	103.74967
## 1089	95.11	89.41441	100.80559	86.39934	103.82066
## 1090	95.11	89.36837	100.85163	86.32893	103.89107
## 1091	95.11	89.32269	100.89731	86.25908	103.96092
## 1092	95.11	89.27738	100.94262	86.18977	104.03023
## 1093	95.11	89.23241	100.98759	86.12100	104.09900
## 1094	95.11	89.18778	101.03222	86.05275	104.16725
## 1095	95.11	89.14349	101.07651	85.98501	104.23499
## 1096	95.11	89.09953	101.12047	85.91777	104.30223
## 1097	95.11	89.05588	101.16412	85.85102	104.36898
## 1098	95.11	89.01255	101.20745	85.78475	104.43525
## 1099	95.11	88.96952	101.25048	85.71894	104.50106
## 1100	95.11	88.92679	101.29321	85.65360	104.56640
## 1101	95.11	88.88435	101.33565	85.58870	104.63130
## 1102	95.11	88.84221	101.37779	85.52424	104.69576
## 1103	95.11	88.80034	101.41966	85.46021	104.75979
## 1104	95.11	88.75875	101.46125	85.39660	104.82340
## 1105	95.11	88.71743	101.50257	85.33341	104.88659
## 1106	95.11	88.67637	101.54363	85.27062	104.94938
## 1107	95.11	88.63558	101.58442	85.20823	105.01177
## 1108	95.11	88.59504	101.62496	85.14623	105.07377
## 1109	95.11	88.55475	101.66525	85.08461	105.13539
## 1110	95.11	88.51471	101.70529	85.02338	105.19662
## 1111	95.11	88.47491	101.74509	84.96251	105.25749
## 1112	95.11	88.43535	101.78465	84.90200	105.31800
## 1113	95.11	88.39602	101.82398	84.84185	105.37815
## 1114	95.11	88.35692	101.86308	84.78205	105.43795

The plot above shows the expected values for rate of fuel for next 90 observations where the rate of fuel prices seem to be constant for the next observations since the prices have been showing no variation towards the end data.